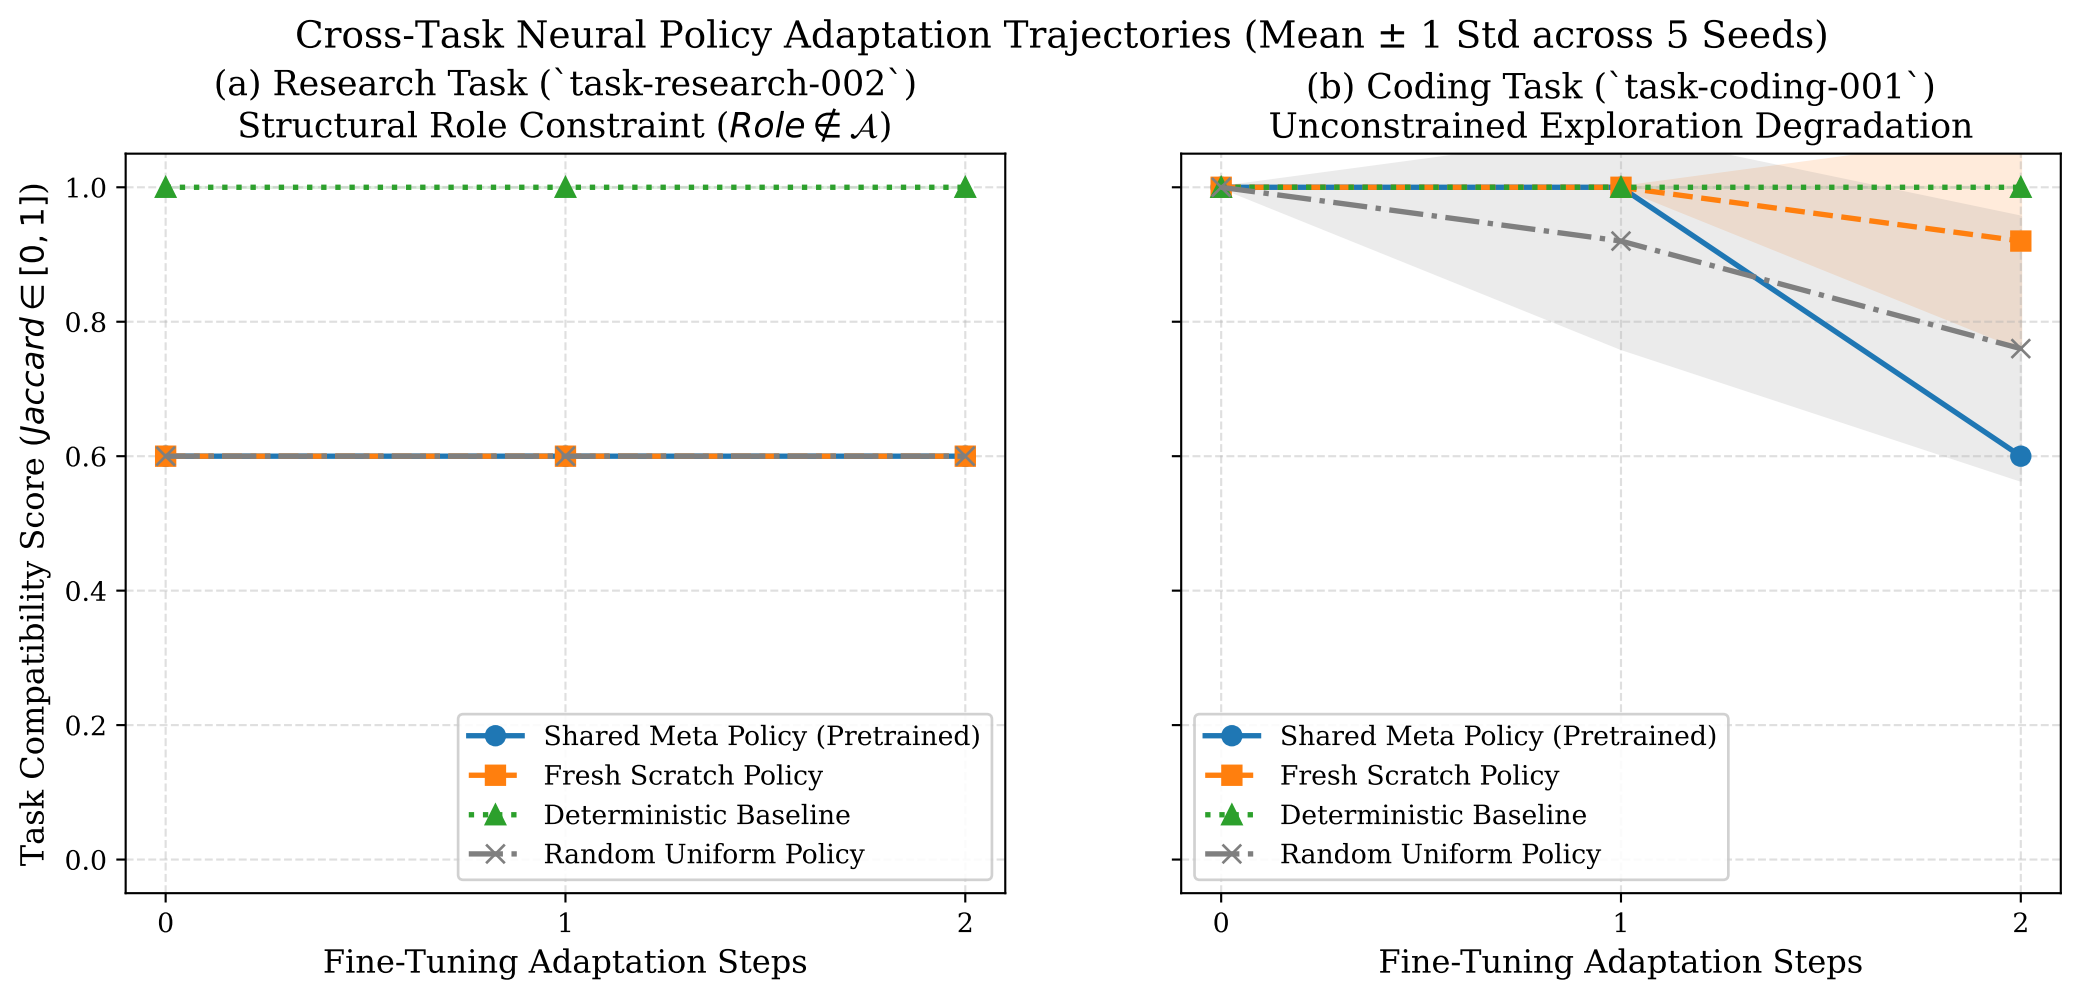
Cross-Task Neural Policy Adaptation Trajectories (Mean ± 1 Std across 5 Seeds)
(a) Research Task (`task-research-002`)
Structural Role Constraint ($Role \notin \mathcal{A}$)
(b) Coding Task (`task-coding-001`)
Unconstrained Exploration Degradation
Task Compatibility Score ($Jaccard \in [0, 1]$)
1.0
0.8
0.6
0.4
0.2
0.0
0
1
2
Fine-Tuning Adaptation Steps
Shared Meta Policy (Pretrained)
Fresh Scratch Policy
Deterministic Baseline
Random Uniform Policy