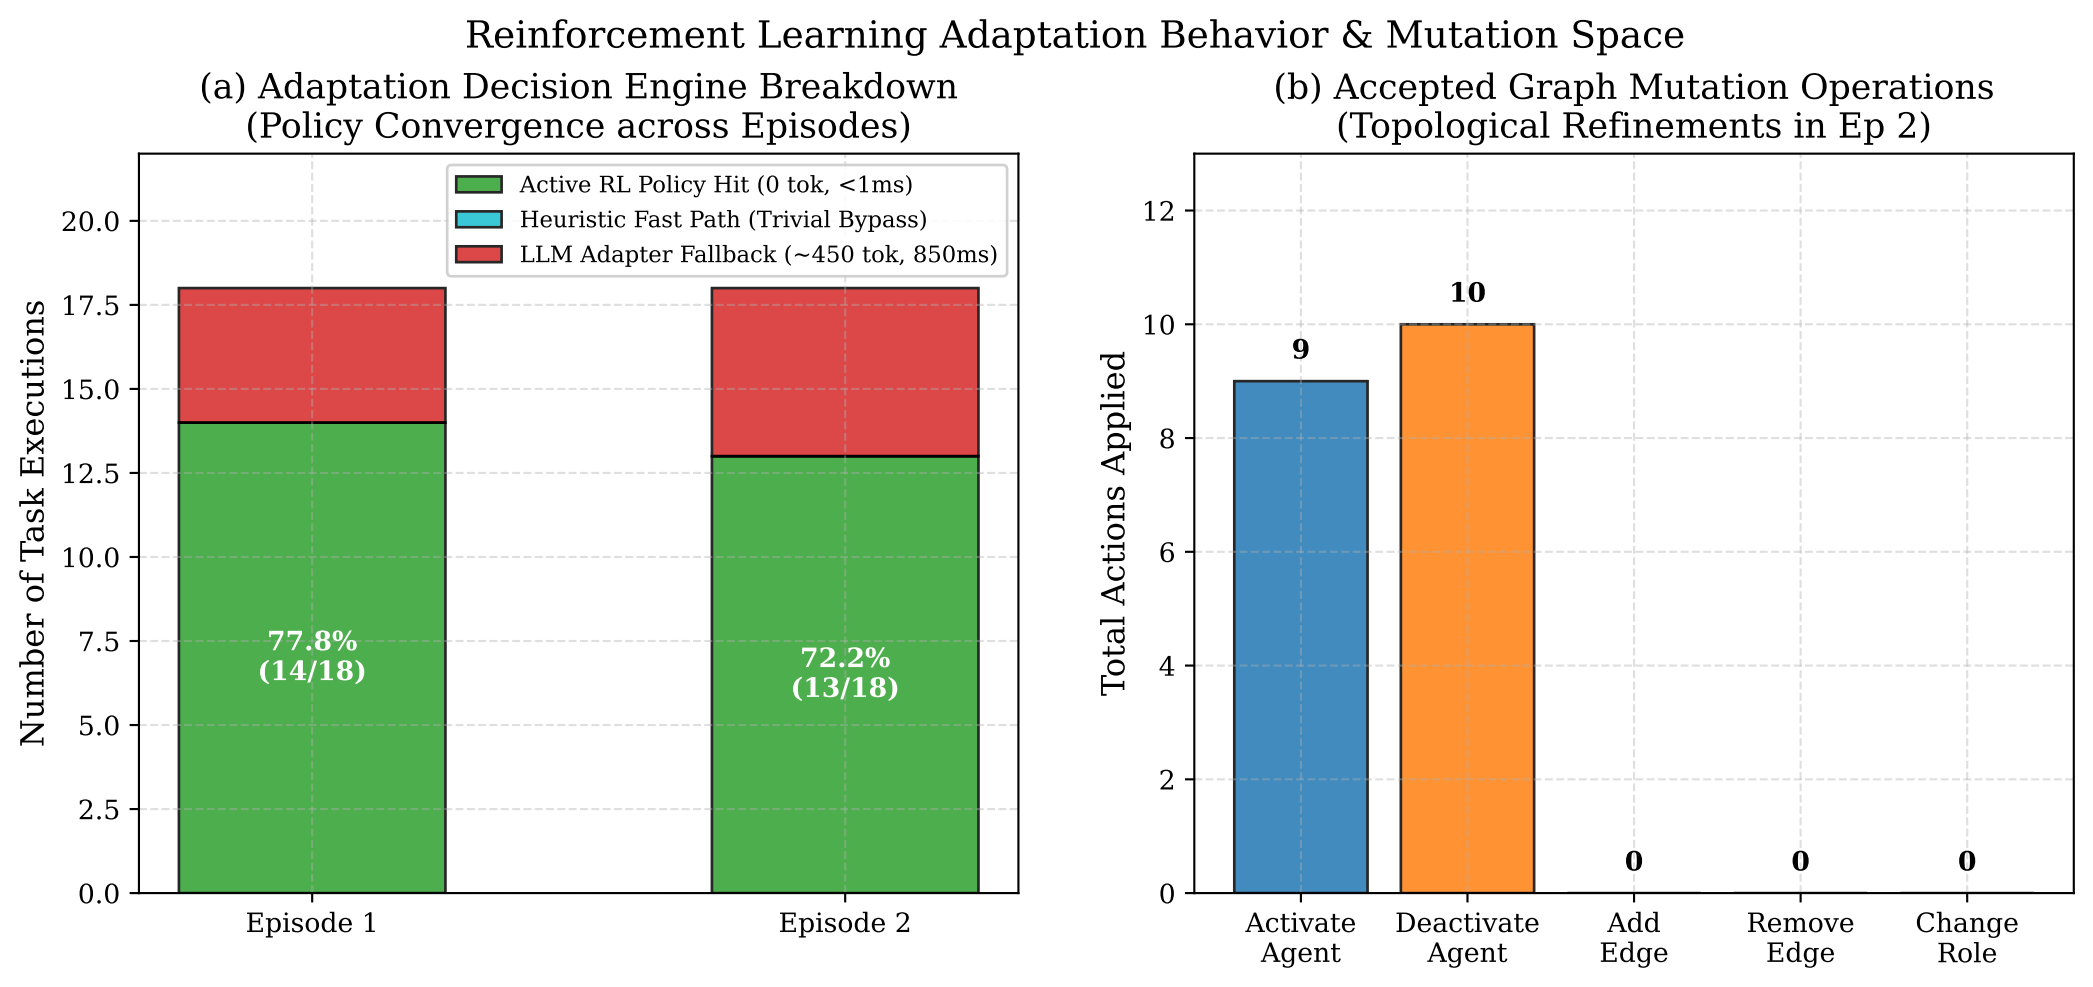
Reinforcement Learning Adaptation Behavior & Mutation Space
(a) Adaptation Decision Engine Breakdown
(Policy Convergence across Episodes)
Active RL Policy Hit (0 tok, <1ms)
Heuristic Fast Path (Trivial Bypass)
LLM Adapter Fallback (~450 tok, 850ms)
Number of Task Executions
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
20.0
77.8%
(14/18)
72.2%
(13/18)
Episode 1
Episode 2
(b) Accepted Graph Mutation Operations
(Topological Refinements in Ep 2)
Total Actions Applied
0
2
4
6
8
10
12
9
10
0
0
0
Activate Agent
Deactivate Agent
Add Edge
Remove Edge
Change Role